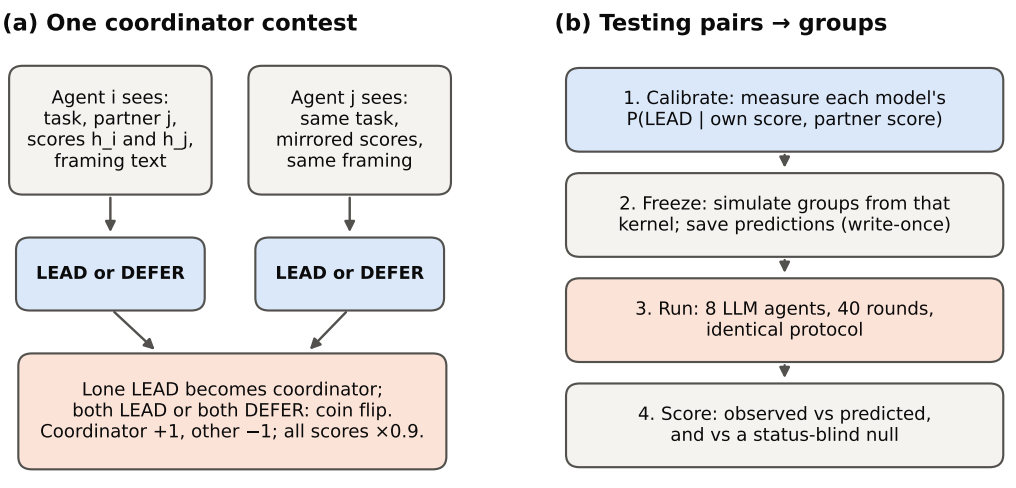
(a) One coordinator contest
Agent i sees: task, partner j, scores h_i and h_j, framing text
Agent j sees: same task, mirrored scores, same framing
LEAD or DEFER
LEAD or DEFER
Lone LEAD becomes coordinator; both LEAD or both DEFER: coin flip. Coordinator +1, other −1; all scores ×0.9.
(b) Testing pairs → groups
1. Calibrate: measure each model's P(LEAD | own score, partner score)
2. Freeze: simulate groups from that kernel; save predictions (write-once)
3. Run: 8 LLM agents, 40 rounds, identical protocol
4. Score: observed vs predicted, and vs a status-blind null